

1 Prerequisite Definitions

Alphabets Σ , and Γ are finite nonempty sets of symbols.

A *string* is a finite sequence of zero or more symbols from an alphabet.

Σ^* is the set of all strings over alphabet Σ .

ε is the empty string and cannot be in Σ .

A *problem* is a mapping from strings to strings.

A *decision problem* is a problem whose output is yes/no (or often accept/reject).

A decision problem be thought of as the set of all strings for which the function outputs “accept”.

A *language* is a set of strings, so any set $S \subseteq \Sigma^*$ is a language, even \emptyset . Thus, decision problems are equivalent to languages.

2 Regular Languages

$L(M)$ is the language accepted by machine M .

A deterministic finite automaton is a 5-tuple $M = (Q, \Sigma, \delta, q_0, F)$, where

- Q is a finite set of states,
- Σ is an alphabet,
- $\delta : Q \times \Sigma \rightarrow Q$ is a transition function describing its transitions and labels,
- $q_0 \in Q$ is the starting state, and
- $F \subseteq Q$ is a set of accepting states.

If δ is not fully specified, we assume an implicit transition to an *error state*.

A deterministic finite automaton M accepts input string $w = w_1 w_2 \dots w_n$ ($w_i \in \Sigma^*$) if there exists a sequence of states $r_0, r_1, r_2, \dots, r_n$ ($r_i \in Q$) such that

- $r_0 = q_0$,
- for all $i \in \{1, \dots, n\}$, $r_i = \delta(r_{i-1}, w_i)$, and
- $r_n \in F$.

$r_0, r_1, r_2, \dots, r_n$ are the sequence of states visited during the machine's computation.

A non-deterministic finite automaton is a 5-tuple $M = (Q, \Sigma, \delta, q_0, F)$, where

- Q, Σ, q_0, F are the same as a deterministic finite automaton's, and
- $\delta : Q \times (\Sigma \cup \{\varepsilon\}) \rightarrow 2^Q$.

A non-deterministic finite automaton accepts the string $w = w_1 w_2 \dots w_n$ ($w_i \in \Sigma^*$) if there exist a string $y = y_1 y_2 \dots y_m$ ($y_i \in (\Sigma \cup \{\varepsilon\})^*$) and a sequence $r = r_0, r_1, \dots, r_n$ ($r_i \in Q$) such that

- $w = y_1 \circ y_2 \circ \dots \circ y_m$ (i.e. y is w with some ε inserted),
- $r_0 = q_0$,
- for all $i = \{1, \dots, m\}$, $r_i \in \delta(r_{i-1}, q_i)$, and
- $r_m \in F$.

The ε -closure for any set $S \subseteq Q$ is denoted $E(S)$, which is the set of all states in Q that can be reachable by following any number of ε -transition.

Theorem 1. A non-deterministic finite automaton can be converted to an equivalent deterministic finite automaton.

A *regular language* is any language accepted by some finite automaton. The set of all regular languages is called the *class of regular languages*.

Theorem 2. Regular languages are closed under

- Concatenation $L_1 \circ L_2 = \{x \circ y : x \in L_1 \text{ and } y \in L_2\}$. Note: $L_1 \not\subseteq L_1 \circ L_2$.
- Union $L_1 \cup L_2 = \{x : x \in L_1 \text{ or } x \in L_2\}$.
- Intersection $L_1 \cap L_2 = \{x : x \in L_1 \text{ and } x \in L_2\}$.
- Complement $\bar{L} = \Sigma^* \setminus L = \{x : x \notin L\}$.
- Star $L^* = \{x_1 \circ x_2 \circ \dots \circ x_k : x_i \in L \text{ and } k \geq 0\}$.

R is a regular expression if R is

- $a \in \Sigma$,
- ε ,
- \emptyset ,
- $R_1 \cup R_2$, or $R_1 | R_2$,
- $R_1 \circ R_2$, or $R_1 R_2$,
- R_1^* ,
- Shorthand: $\Sigma = (a_1 | a_2 | \dots | a_k)$, $a_i \in \Sigma$,

where R_i is a regular expression.

Identities of Regular Languages

- $\emptyset \cup R = R \cup \emptyset = R$
- $\emptyset \circ R = R \circ \emptyset = \emptyset$
- $\varepsilon \circ R = R \circ \varepsilon = R$
- $\varepsilon^* = \varepsilon$
- $\emptyset^* = \emptyset$
- $\emptyset \cup R \circ R^* = R \circ R^* \cup \varepsilon = R^*$
- $(a|b)^* = (a^*|b^*)^* = (a^*b^*)^* = (a^*|b)^* = (a|b^*)^* = a^*(ba^*)^* = b^*(ab^*)^*$

Theorem 3. Languages accepted by DFAs = languages accepted by NFAs = regular languages

Theorem 4. If L is a finite language, L is regular.

If a computation path of any finite automaton is longer than the number of states it has, there must be a cycle in that computation path.

Lemma 1 (Pumping Lemma). Every regular language satisfies the pumping condition.

Pumping condition: There exists an integer p such that for every string $w \in L$, with $|w| \geq p$, there exist strings $x, y, z \in \Sigma^*$ with $w = xyz$, $y \neq \varepsilon$, $|xy| \leq p$ such that for all $i \geq 0$, $xy^i z \in L$.

Negation of pumping condition: For all integers p , there exists a string $w \in L$, with $|w| \geq p$, for all $x, y, z \in \Sigma^*$ with $w = xyz$, $y \neq \varepsilon$, $|xy| \leq p$, there exists $i \geq 0, i \neq 1$ such that $xy^i z \notin L$.

Limitations of finite automata:

- Only read input once, left to right.
- Only finite memory.

3 Context-Free Languages

A pushdown automaton is a 6-tuple $M = (Q, \Sigma, \Gamma, \delta, q_0, F)$, where

- Q is a finite set of states,
- Σ is its input alphabet,
- Γ is its stack alphabet,
- $\delta : Q \times (\Sigma \cup \{\epsilon\}) \times (\Gamma \cup \{\epsilon\}) \rightarrow 2^{Q \times (\Gamma \cup \epsilon)}$ is its transition function,
- $q_0 \in Q$ is its starting state, and
- $F \subseteq Q$ is a finite set of accepting states.

Labels: $a, b \rightarrow c$: if input symbol is a , and top of stack is b , pop it and push c . In other words, input symbol read, stack symbol popped \rightarrow stack symbol pushed, e.g. $0, \epsilon \rightarrow \$$.

Suppose u, v, w are strings of variables and terminals, and there is a rule $A \rightarrow w$. From the string uAv , we can obtain uwv . We write $uAv \rightarrow uwv$, and say uAv yields uwv .

If $u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_k$, then $u_1 \rightarrow^* u_k$, or u_1 derives u_k . There must be a finite number of arrows between u_1 and u_k .

Given a grammar G , the language derived by the grammar is $L(G) = \{w \in \Sigma^* : S \rightarrow^* w \text{ and } S \text{ is the start variable}\}$

Context-free grammar: the lhs of rules is a single variable, rhs is any string of variables and terminals. A *context-free language* is one that can be derived from a context-free grammar. An example context-free grammar is $G = (V, \Sigma, R, \langle \text{EXPR} \rangle)$, where $V = \{\langle \text{EXPR} \rangle, \langle \text{TERM} \rangle, \langle \text{FACTOR} \rangle\}$,

$\Sigma = \{a, +, \times, (,)\}$, and $R = \{\langle \text{EXPR} \rangle \rightarrow \langle \text{EXPR} \rangle + \langle \text{TERM} \rangle, \langle \text{TERM} \rangle \rightarrow \langle \text{TERM} \rangle \times \langle \text{FACTOR} \rangle, \langle \text{FACTOR} \rangle \rightarrow \langle \text{EXPR} \rangle\}$.

A *left-most derivation* is a sequence $S \rightarrow u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_k \rightarrow w$ where each step applies a rule to the left-most variable. A grammar is *ambiguous* when it has multiple left-most derivations for the same string.

Theorem 5. A language L is recognized by a pushdown automaton iff L is described by a context-free grammar.

Theorem 6. Context-free languages are closed under union, concatenation, star.

4 Recognizable Languages

Differences from previous models

- The input is written on tape.
- It can write to the tape.
- It can move left and right on tape.
- It halts immediately when it reaches an accepting or rejecting state. The rejecting state must exist but may not be shown.

A deterministic Turing machine is a 7-tuple $M = (Q, \Sigma, \Gamma, \delta, q_0, q_{\text{accept}}, q_{\text{reject}})$, where

- Q is its finite non-empty set of states,
- Σ is its input alphabet,
- Γ is its tape alphabet ($\Sigma \subset \Gamma$ and $\sqcup \in \Gamma \setminus \Sigma$),

- $\delta : Q \times \Gamma \rightarrow Q \times \Gamma \times \{L, R\}$ is its transition function,
- $q_0 \in Q$ is its starting state,
- $q_{\text{accept}} \in Q$ is its accepting state, and
- $q_{\text{reject}} \in Q$ is its rejecting state ($q_{\text{reject}} \neq q_{\text{accept}}$).

Labels: $a \rightarrow b, R$: if tape symbol is a , write b and move head right. $a \rightarrow R$: if tape symbol is a , move head right. $a, b, c \rightarrow R$: if tape symbol is a, b , or c , move head right.

On input x , a Turing machine can (1) accept, (2) reject, or (3) run in an infinite loop.

The language *recognized* by a Turing machine M is $L(M) = \{x : \text{on input } x, M \text{ halts in } q_{\text{accept}}\}$. A language is *recognizable* if there exists a Turing machine which recognizes it.

Regular languages \subseteq context-free languages \subseteq decidable languages \subseteq recognizable languages

A *configuration* is a way to describe the entire state of the Turing machine. It is a string aqb where $a \in \Gamma^*, q \in Q, b \in \Gamma^*$, which indicates that q is the current state of the Turing machine, the tape content currently is ab and its head is currently pointing at the first symbol of b . Any Turing machine halts if its configuration is of the form $aq_{\text{accept}}b$, or $aq_{\text{reject}}b$ for any ab . $\text{Config}(i)$ uniquely determines $\text{Config}(i+1)$.

Theorem 7. Every k -tape Turing machine has an equivalent single tape Turing machine.

If the alphabet of the multitape Turing machine is Γ , we can make the single tape Turing machine's alphabet $(\Gamma \cup \{\#\}) \times \{\text{normal}, \text{bold}\}$.

A non-deterministic Turing machine is a 7-tuple $M = (Q, \Sigma, \Gamma, \delta, q_0, q_{\text{accept}}, q_{\text{reject}})$, where the only difference from a deterministic Turing machine is the transition function $\delta : Q \times \Gamma \rightarrow 2^{Q \times \Gamma \times \{L, R\}}$.

A non-deterministic Turing machine accepts its input iff some node in the configuration tree has q_{accept} . It does not accept its input iff the configuration tree grows forever (infinite loop) or no node in the tree has q_{accept} .

Acceptance of a non-deterministic Turing machine: input w is accepted if there exist configurations c_0, c_1, \dots, c_k where

- $c_0 = q_{\text{start}}w$, and
- $c_i \Rightarrow c_{i+1}$ (c_{i+1} is a possible configuration from c_i , following the transition function δ).

The outcomes could be

- w is accepted, i.e. there exists a node in the tree which is an accepting configuration,
- w is explicitly rejected, i.e. the tree is finite but no node is an accepting configuration (all leaves are rejecting configurations), or
- the non-deterministic Turing machine runs forever on w , i.e. the tree is infinite but no node is an accepting configuration (there might be finite branches terminating in a rejecting configuration in the tree).

A Turing machine is a *decider* if it halts on all inputs, i.e. it either rejects or accepts all inputs.

Theorem 8. *Every non-deterministic Turing machine has an equivalent deterministic Turing machine. If that non-deterministic Turing machine is a decider, there is an equivalent deterministic Turing machine decider.*

Theorem 9. *Recognizable languages are closed under union, intersection, concatenation, star.*

Implementation level description of a multitape Turing machine for $L = \{x\#x : x \in \{0, 1\}^*\}$:

- Scan the first head to the right until it reads a #. Move right. The second head is still at the start of the second tape.
- Repeatedly read symbol from the first tape (reject if the symbol is not 0 or 1), write it to the second tape, and move both heads right, until seeing a blank on the first tape.
- Move the first head left until a # is under it. Replace the symbol with a blank ($_$).
- Move both heads left until they reach the start of their respective

tapes (using the \$ sign hack to mark the start of the tape).

- Repeat until seeing a blank on both tapes.
 - If the symbols on the two tapes differ, reject.
 - Otherwise, move both head right.

$\langle O \rangle$ is a string encoding for the object O .

Cardinality of Sets: two sets A and B have the same *cardinality* if there exists a bijection $f : A \rightarrow B$.

$\mathbb{N} = \{1, 2, 3, \dots\}$ is the set of all natural numbers. A set is *finite* if it has

a bijection to $\{1..n\}$ for some natural number n . A set is *countably infinite* if it has the same cardinality as \mathbb{N} . A set is *countable* or *at most countable* if it is finite or countably infinite.

Lemma 2. *Any language L is countable.*

Lemma 3. *The set of all Turing machines is countable.*

Lemma 4. *The set \mathcal{B} of all infinite bit-sequences is not countable.*

Lemma 5. 2^{Σ^*} is uncountable.

$A_{TM} = \{\langle M, w \rangle : M \text{ accepts } w\}$ is recognizable but not decidable.